

Lecture 11: Logistic Regression and Classification

In classification problem, the response variables y are discrete, representing different categories. For simplicity, we will first introduce the **binary classification case** -- y has only two categories, denoted as 0 and 1.

Model-setup

Assumption 1: Dependent on the variable x , the response variable y has different **probabilities** to take value in 0 or 1. Instead of predicting exact value of 0 or 1, we are actually predicting the **probabilities**.

Assumption 2: Logistic function assumption. Given x , what is the probability to observe $y = 1$?

$$P(y = 1 | \mathbf{x}) = f(\mathbf{x}; \beta) = \frac{1}{1 + \exp(-\tilde{x}\beta)} =: \sigma(\tilde{x}\beta).$$

where $\sigma(z) = \frac{1}{1 + \exp(-z)}$ is called [standard logistic function](https://en.wikipedia.org/wiki/Logistic_regression#:~:text=Logistic%20regression%20is%20a%20statistical,a%20form%20of%20b)

(https://en.wikipedia.org/wiki/Logistic_regression#:~:text=Logistic%20regression%20is%20a%20statistical,a%20form%20of%20b) or sigmoid function in deep learning. Recall that $\beta \in \mathbb{R}^{p+1}$ and \tilde{x} is the "augmented" sample with first element one to incorporate intercept in the linear function.

Equivalent expression:

- Denote $p = P(y = 1 | \mathbf{x})$, then we can write in linear form

$$\ln \frac{p}{1-p} = \tilde{x}\beta$$

- Since y only takes value in 0 or 1, we have

$$P(y | \mathbf{x}, \beta) = f(\mathbf{x}; \beta)^y (1 - f(\mathbf{x}; \beta))^{1-y}$$

MLE (Maximum Likelihood Estimation)

Assume the samples are independent. The overall probability to witness the whole training dataset

$$\begin{aligned} P(\mathbf{y} | \mathbf{X}; \beta) \\ &= \prod_{i=1}^N P(y^{(i)} | \mathbf{x}^{(i)}; \beta) \\ &= \prod_{i=1}^N f(\mathbf{x}^{(i)}; \beta)^{y^{(i)}} (1 - f(\mathbf{x}^{(i)}; \beta))^{(1-y^{(i)})}. \end{aligned}$$

By maximizing the logarithm of likelihood function, then we derive the **loss function** to be minimized $L(\beta) = L(\beta; \mathbf{X}, \mathbf{y}) = -\frac{1}{N} \sum_{i=1}^N \ln \big(f(\mathbf{x}^{(i)}; \beta)^{y^{(i)}} (1 - f(\mathbf{x}^{(i)}; \beta))^{(1-y^{(i)})} \big)$

- $(1 - y^{(i)}) \ln (1 - f(\mathbf{x}^{(i)}; \beta))$.

The loss function also has clear probabilistic interpretations. Given i -th sample, the vector of true labels $(y^i, 1 - y^i)$ can also be viewed as the probability distribution. Then the loss function is the mean of all [cross entropy](https://en.wikipedia.org/wiki/Cross_entropy) (https://en.wikipedia.org/wiki/Cross_entropy) across samples, i.e. "distance" between true probability distribution and estimated probability distribution via logistic model.

Algorithm

Take the gradient (left as exercise -- if you like)

$$\frac{\partial L(\beta)}{\partial \beta_k} = \frac{1}{N} \sum_{i=1}^N (\sigma(\tilde{x}^{(i)} \beta) - y^{(i)}) \tilde{x}_k^{(i)}.$$

In vector form

$$\nabla_{\beta} (L(\beta)) = \sum_{i=1}^N (\sigma(\tilde{x}^{(i)}) - y^{(i)}) \tilde{x}^{(i)} = \frac{1}{N} \sum_{i=1}^N (f(\mathbf{x}^{(i)}; \beta) - y^{(i)}) \tilde{x}^{(i)}.$$

This is still the nonlinear function of β , indicating that we cannot derive something like "normal equations" in OLS. The solution here is [numerical optimization \(https://github.com/Jaewan-Yun/optimizer-visualization\)](https://github.com/Jaewan-Yun/optimizer-visualization).

The simplest algorithm in optimization is [gradient descent \(GD\)](https://en.wikipedia.org/wiki/Gradient_descent#:~:text=Gradient%20descent%20is%20a%20first,function%20at%20the%20curre)

(https://en.wikipedia.org/wiki/Gradient_descent#:~:text=Gradient%20descent%20is%20a%20first,function%20at%20the%20curre

$$\beta^{k+1} = \beta^k - \eta \nabla L(\beta^k).$$

Here the step size η is also called **learning rate** in machine learning. Note that it is indeed the Euler's scheme to solve the ODE

$$\dot{\beta} = -\nabla L(\beta).$$

By setting certain stopping criterion for GD, we think that we have approximated the optimized solution $\hat{\beta}$.

Making predictions

Now with the estimated $\hat{\beta}$ and given a new data x^{new} , we calculate the probability that $y^{new} = 1$ as $f(\mathbf{x}; \beta)$. If is greater than 0.5, we assign that $y^{new} = 1$.

For the whole dataset, the **accuracy** is defined as ratio of number of correct predictions to the total number of samples.

Example Code

```

In [1]: import numpy as np

class myLogisticRegression_binary():
    """ Logistic Regression classifier -- this only works for the binary case.
    Parameters:
    -----
    learning_rate: float
        The step length that will be taken when following the negative gradient during
        training.
    """
    def __init__(self, learning_rate=.1):

        # learning rate can also be in the fit method
        self.learning_rate = learning_rate

    def fit(self, data, y, n_iterations = 1000):
        """
        don't forget the document string in methods
        """
        ones = np.ones((data.shape[0],1)) # column of ones
        X = np.concatenate((ones, data), axis = 1) # the augmented matrix, \tilde{X}
        in our lecture
        eta = self.learning_rate

        beta = np.zeros(np.shape(X)[1]) # initialize beta, can be other choices

        for k in range(n_iterations):
            dbeta = self.loss_gradient(beta,X,y) # write another function to compute
            gradient
            beta = beta - eta * dbeta # the formula of GD
            # this step is optional -- just for inspection purposes
            if k % 500 == 0: # pprint loss every 500 steps
                print("loss after", k+1, "iterations is: ", self.loss(beta,X,y))

            self.coeff = beta

    def predict(self, data):
        ones = np.ones((data.shape[0],1)) # column of ones
        X = np.concatenate((ones, data), axis = 1) # the augmented matrix, \tilde{X}
        in our lecture
        beta = self.coeff # the estimated beta
        y_pred = np.round(self.sigmoid(np.dot(X,beta))).astype(int) # >0.5: ->1 else,
        ->0 -- note that we always use Numpy universal functions when possible
        return y_pred

    def score(self, data, y_true):
        ones = np.ones((data.shape[0],1)) # column of ones
        X = np.concatenate((ones, data), axis = 1) # the augmented matrix, \tilde{X}
        in our lecture
        y_pred = self.predict(data)
        acc = np.mean(y_pred == y_true) # number of correct predictions/N
        return acc

    def sigmoid(self, z):
        return 1.0 / (1.0 + np.exp(-z))

    def loss(self,beta,X,y):
        f_value = self.sigmoid(np.matmul(X,beta))
        loss_value = np.log(f_value + 1e-10) * y + (1.0 - y)* np.log(1 - f_value + 1e
        -10) # avoid nan issues
        return -np.mean(loss_value)

    def loss_gradient(self,beta,X,y):
        f_value = self.sigmoid(np.matmul(X,beta))
        gradient_value = (f_value - y).reshape(-1,1)*X # this is the hardest expressi
        on -- check yourself. It's called Numpy broadcasting
        return np.mean(gradient_value, axis=0)

```

```
In [2]: from sklearn.datasets import load_breast_cancer
X, y = load_breast_cancer(return_X_y = True)
```

```
In [3]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1, random_state=42)
```

```
In [4]: X_train.shape
```

```
Out[4]: (512, 30)
```

```
In [5]: %%time
lg = myLogisticRegression_binary(learning_rate=1e-5)
lg.fit(X_train,y_train,n_iterations = 20000) # what about increase n_iterations?
```

```
loss after 1 iterations is: 0.7704000919325609
loss after 501 iterations is: 0.3038878556607375
loss after 1001 iterations is: 0.2646726705164665
loss after 1501 iterations is: 0.24813479245950681
loss after 2001 iterations is: 0.2389480595788275
loss after 2501 iterations is: 0.23311785521728876
loss after 3001 iterations is: 0.22909746536348718
loss after 3501 iterations is: 0.22615149966747047
loss after 4001 iterations is: 0.2238860140178029
loss after 4501 iterations is: 0.22207285161370108
loss after 5001 iterations is: 0.22057221113785216
loss after 5501 iterations is: 0.2192946215702638
loss after 6001 iterations is: 0.21818078310948252
loss after 6501 iterations is: 0.21719023774728446
loss after 7001 iterations is: 0.21629469731306186
loss after 7501 iterations is: 0.21547395937965436
loss after 8001 iterations is: 0.21471332436186458
loss after 8501 iterations is: 0.21400191582326006
loss after 9001 iterations is: 0.21333156155309688
loss after 9501 iterations is: 0.21269603244872282
loss after 10001 iterations is: 0.21209051523044697
loss after 10501 iterations is: 0.21151124122819925
loss after 11001 iterations is: 0.21095522130259972
loss after 11501 iterations is: 0.21042005414951925
loss after 12001 iterations is: 0.20990378610015575
loss after 12501 iterations is: 0.20940480753853602
loss after 13001 iterations is: 0.2089217756675304
loss after 13501 iterations is: 0.20845355643721925
loss after 14001 iterations is: 0.20799918054355343
loss after 14501 iterations is: 0.20755780984807357
loss after 15001 iterations is: 0.2071287115765159
loss after 15501 iterations is: 0.2067112383654315
loss after 16001 iterations is: 0.20630481273382617
loss after 16501 iterations is: 0.20590891492308783
loss after 17001 iterations is: 0.20552307331501238
loss after 17501 iterations is: 0.20514685683332629
loss after 18001 iterations is: 0.20477986887872715
loss after 18501 iterations is: 0.20442174245513262
loss after 19001 iterations is: 0.2040721362254978
loss after 19501 iterations is: 0.20373073129634595
CPU times: user 1.48 s, sys: 50.3 ms, total: 1.53 s
Wall time: 1.53 s
```

```
In [6]: lg.score(X_test,y_test)
```

```
Out[6]: 1.0
```

```
In [7]: lg.score(X_train,y_train)
```

```
Out[7]: 0.916015625
```

```
In [8]: lg.predict(X_test)
```

```
Out[8]: array([[1, 0, 0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 1,
               0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1,
               1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1]])
```

```
In [9]: y_test
```

```
Out[9]: array([[1, 0, 0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 0, 0, 1,
               0, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 1,
               1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 1, 1, 1]])
```

```
In [10]: lg.coef
```

```
Out[10]: array([ 2.66882874e-03,  1.83927907e-02, -4.36239390e-03,  8.66487934e-02,
                1.49727981e-02,  5.50578751e-05, -6.71061493e-04, -1.14024608e-03,
               -4.51040831e-04,  1.06678514e-04,  8.62208844e-05,  1.72299429e-04,
                4.51211649e-04, -2.32558967e-03, -2.93966958e-02, -5.28628701e-06,
               -1.83325756e-04, -2.46370803e-04, -5.80574855e-05, -9.52561495e-06,
               -1.16457037e-05,  1.92488199e-02, -1.67105144e-02,  6.60427872e-02,
               -2.88220491e-02, -2.09664782e-05, -2.48568195e-03, -3.30713576e-03,
               -8.60537728e-04, -1.96407733e-04, -8.95680417e-05])
```

```
In [11]: from sklearn.linear_model import LogisticRegression
         clf = LogisticRegression(random_state=0)
         clf.fit(X_train,y_train)
         clf.score(X_test,y_test)
```

```
/Users/cliffzhou/opt/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_log
istic.py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)

```
Out[11]: 0.9824561403508771
```

```
In [12]: clf.score(X_train,y_train)
```

```
Out[12]: 0.953125
```

It's very normal that our result is different with sklearn. In sklearn [logistic regression \(https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html#sklearn.linear_model.LogisticRegression\)](https://scikit-learn.org/stable/modules/generated/sklearn.linear_model.LogisticRegression.html#sklearn.linear_model.LogisticRegression), they use different algorithms to solve the optimization problem, and even the model is different (adding regularization terms!)

Multi-class Case

Model

Let $\tilde{x} \in \mathbb{R}^{p+1}$ denotes the augmented row vector (one sample). We approximate the probabilities to take value in K classes as

$$f(\mathbf{x}; W) = \begin{pmatrix} P(y = 1|\mathbf{x}; \mathbf{w}) \\ P(y = 2|\mathbf{x}; \mathbf{w}) \\ \vdots \\ P(y = K|\mathbf{x}; \mathbf{w}) \end{pmatrix} = \frac{1}{\sum_{k=1}^K \exp(\tilde{x}\mathbf{w}_k)} \begin{pmatrix} \exp(\tilde{x}\mathbf{w}_1) \\ \exp(\tilde{x}\mathbf{w}_2) \\ \vdots \\ \exp(\tilde{x}\mathbf{w}_K) \end{pmatrix}.$$

where we have K sets of parameters, $\mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_K$, and the sum factor normalizes the results to be a probability.

\mathbf{W} is an $(p+1) \times K$ matrix containing all K sets of parameters, obtained by concatenating $\mathbf{w}_1, \mathbf{w}_2, \dots, \mathbf{w}_K$ into columns, so that $\mathbf{w}_k = (w_{k0}, \dots, w_{kp})^\top \in \mathbb{R}^{p+1}$.

$$\mathbf{W} = \begin{pmatrix} | & | & | & | \\ \mathbf{w}_1 & \mathbf{w}_2 & \dots & \mathbf{w}_K \\ | & | & | & | \end{pmatrix},$$

and $\tilde{X}\mathbf{W}$ is valid and useful in vectorized code.

Loss function

Define the following indicator function:

$$1_{\{y=k\}} = 1_{\{k\}}(y) = \delta_{yk} = \begin{cases} 1 & \text{when } y = k, \\ 0 & \text{otherwise.} \end{cases}$$

Loss function is again using the cross entropy:

$$\begin{aligned} L(\mathbf{W}; X, \mathbf{y}) &= -\frac{1}{N} \sum_{i=1}^N \sum_{k=1}^K \left\{ 1_{\{y^{(i)}=k\}} \ln P(y^{(i)} = k|\mathbf{x}^{(i)}; \mathbf{w}) \right\} \\ &= -\frac{1}{N} \sum_{i=1}^N \sum_{k=1}^K \left\{ 1_{\{y^{(i)}=k\}} \ln \left(\frac{\exp(\tilde{x}^{(i)} \mathbf{w}_k)}{\sum_{m=1}^K \exp(\tilde{x}^{(i)} \mathbf{w}_m)} \right) \right\}. \end{aligned}$$

Notice that for each term in the summation over N (i.e. fix sample i), only one term is non-zero in the sum of K elements due to the indicator function.

Gradient descent

After **careful calculation**, the gradient of L with respect the whole k -th set of weights is then:

$$\frac{\partial L}{\partial \mathbf{w}_k} = \frac{1}{N} \sum_{i=1}^N \left(\frac{\exp(\tilde{x}^{(i)} \mathbf{w}_k)}{\sum_{m=1}^K \exp(\tilde{x}^{(i)} \mathbf{w}_m)} - 1_{\{y^{(i)}=k\}} \right) \tilde{x}^{(i)} \in \mathbb{R}^{p+1}.$$

In writing the code, it's helpful to make this as the column vector, and stack all the K gradients together as a new matrix $\mathbf{dW} \in \mathbb{R}^{(p+1) \times K}$. This makes the update of matrix \mathbf{W} very convenient in gradient descent.

Prediction

The largest estimated probability's class as this sample's predicted label.

$$\hat{y} = \arg \max_j P(y = j|\mathbf{x}),$$

```

In [16]: import numpy as np

class myLogisticRegression():
    """ Logistic Regression classifier -- this also works for the multiclass case.
    Parameters:
    -----
    learning_rate: float
        The step length that will be taken when following the negative gradient during
        training.
    """
    def __init__(self, learning_rate=.1):

        # learning rate can also be in the fit method
        self.learning_rate = learning_rate

    def fit(self, data, y, n_iterations = 1000):
        """
        don't forget the document string in methods
        """
        self.K = max(y)+1 # specify number of classes in y
        ones = np.ones((data.shape[0],1)) # column of ones
        X = np.concatenate((ones, data), axis = 1) # the augmented matrix, \tilde{X}
        in our lecture
        eta = self.learning_rate

        W = np.zeros((np.shape(X)[1],max(y)+1)) # initialize beta, can be other choices

        for k in range(n_iterations):
            dW = self.loss_gradient(W,X,y) # write another function to compute gradient
            nt
            W = W - eta * dW # the formula of GD
            # this step is optional -- just for inspection purposes
            if k % 500 == 0: # print loss every 500 steps
                print("loss after", k+1, "iterations is: ", self.loss(W,X,y))

            self.coeff = W

    def predict(self, data):
        ones = np.ones((data.shape[0],1)) # column of ones
        X = np.concatenate((ones, data), axis = 1) # the augmented matrix, \tilde{X}
        in our lecture
        W = self.coeff # the estimated W
        y_pred = np.argmax(self.sigma(X,W), axis =1) # the category with largest probability
        ability
        return y_pred

    def score(self, data, y_true):
        ones = np.ones((data.shape[0],1)) # column of ones
        X = np.concatenate((ones, data), axis = 1) # the augmented matrix, \tilde{X}
        in our lecture
        y_pred = self.predict(data)
        acc = np.mean(y_pred == y_true) # number of correct predictions/N
        return acc

    def sigma(self, X, W):
        s = np.exp(np.matmul(X,W))
        total = np.sum(s, axis=1).reshape(-1,1)
        return s/total

    def loss(self, W, X, y):
        f_value = self.sigma(X,W)
        K = self.K
        loss_vector = np.zeros(X.shape[0])
        for k in range(K):
            loss_vector += np.log(f_value+1e-10)[: ,k] * (y == k) # avoid nan issues
        return -np.mean(loss_vector)

```

```

def loss_gradient(self,W,X,y):
    f_value = self.sigma(X,W)
    K = self.K
    dLdW = np.zeros((X.shape[1],K))
    for k in range(K):
        dLdWk =(f_value[:,k] - (y==k)).reshape(-1,1)*X # Numpy broadcasting
        dLdW[:,k] = np.mean(dLdWk, axis=0) # RHS is 1D Numpy array -- so you ca
n safely put it in the k-th column of 2D array dLdW
    return dLdW

```

```

In [14]: from sklearn.datasets import load_digits
X,y = load_digits(return_X_y = True)
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1, random_state
=42)

```

```

In [17]: lg = myLogisticRegression(learning_rate=1e-4)
lg.fit(X_train,y_train,n_iterations = 20000) # what about change the parameters?

```

```

loss after 1 iterations is:  2.2975031101988965
loss after 501 iterations is:  0.9747646840265886
loss after 1001 iterations is:  0.6271544957404386
loss after 1501 iterations is:  0.48465074291917476
loss after 2001 iterations is:  0.4067886795416971
loss after 2501 iterations is:  0.3569853787369549
loss after 3001 iterations is:  0.3219498860091718
loss after 3501 iterations is:  0.2957112499207807
loss after 4001 iterations is:  0.27517638606506345
loss after 4501 iterations is:  0.2585728459578632
loss after 5001 iterations is:  0.24480630370680928
loss after 5501 iterations is:  0.23316150090969132
loss after 6001 iterations is:  0.22314954388974834
loss after 6501 iterations is:  0.21442400929215735
loss after 7001 iterations is:  0.2067320488693829
loss after 7501 iterations is:  0.19988447822601138
loss after 8001 iterations is:  0.19373672640885967
loss after 8501 iterations is:  0.1881762834198265
loss after 9001 iterations is:  0.1831141858457873
loss after 9501 iterations is:  0.17847909502105724
loss after 10001 iterations is:  0.17421308722718495
loss after 10501 iterations is:  0.17026860268039193
loss after 11001 iterations is:  0.16660619607205016
loss after 11501 iterations is:  0.16319285237565423
loss after 12001 iterations is:  0.16000070825015078
loss after 12501 iterations is:  0.15700606905300005
loss after 13001 iterations is:  0.15418864437799004
loss after 13501 iterations is:  0.1515309472374647
loss after 14001 iterations is:  0.14901781725362154
loss after 14501 iterations is:  0.14663603885543683
loss after 15001 iterations is:  0.14437403299967985
loss after 15501 iterations is:  0.1422216063268606
loss after 16001 iterations is:  0.14016974557619974
loss after 16501 iterations is:  0.13821044795587994
loss after 17001 iterations is:  0.1363365802952386
loss after 17501 iterations is:  0.1345417614013797
loss after 18001 iterations is:  0.13282026324911267
loss after 18501 iterations is:  0.13116692755309206
loss after 19001 iterations is:  0.12957709497826864
loss after 19501 iterations is:  0.12804654479263708

```

```

In [18]: lg.score(X_test,y_test)

```

```

Out[18]: 0.9722222222222222

```

```

In [19]: np.where(lg.predict(X_test)!=y_test)

```

```


Out[19]: (array([ 5, 71, 133, 149, 159]),)

```



```
In [20]: import matplotlib.pyplot as plt
plt.imshow(X_test[149,].reshape(8,8))
```

```
Out[20]: <matplotlib.image.AxesImage at 0x7f9832553b90>
```



```
In [21]: print(lg.predict(X_test)[149],y_test[149])
```

```
5 3
```

For multi-class classification, the [confusion matrix](https://scikit-learn.org/stable/modules/generated/sklearn.metrics.confusion_matrix.html) (https://scikit-learn.org/stable/modules/generated/sklearn.metrics.confusion_matrix.html) can provide as more details.

```
In [22]: from sklearn.metrics import confusion_matrix
confusion_matrix(y_test,lg.predict(X_test))
```

```
Out[22]: array([[17,  0,  0,  0,  0,  0,  0,  0,  0,  0],
 [ 0, 10,  1,  0,  0,  0,  0,  0,  0,  0],
 [ 0,  0, 17,  0,  0,  0,  0,  0,  0,  0],
 [ 0,  0,  0, 16,  0,  1,  0,  0,  0,  0],
 [ 0,  0,  0,  0, 25,  0,  0,  0,  0,  0],
 [ 0,  0,  0,  0,  0, 21,  0,  0,  0,  1],
 [ 0,  0,  0,  0,  0,  0, 19,  0,  0,  0],
 [ 0,  0,  0,  0,  0,  0,  0, 18,  0,  1],
 [ 0,  0,  0,  0,  0,  0,  0,  0,  8,  0],
 [ 0,  0,  0,  0,  0,  0,  0,  0,  1, 24]])
```

Tricks in training: Stochastic Gradient Descent (SGD)

When you're doing the final project, it's very likely that you might lose patience -- training on the 60,000 MNIST data is VERY SLOW! (of course it's not an excuse to abandon the project lol)

To speed up the training process (most importantly the optimization algorithm), there are two directions of general strategies:

- find better algorithm whose convergence is faster (you take less steps to arrive at the minimum)
- save the computational cost within each step

Of course there are trade-offs between these two directions.

Basic observation of SGD: Calculating the gradient in each step is TOO EXPENSIVE!

Recall that in general supervised learning,

$$\nabla_{\beta} L(\beta; X, Y) = \frac{1}{N} \sum_{i=1}^N \nabla_{\beta} l(\beta; x^{(i)}, y^{(i)})$$

It means that we need to implement 60,000 sum calculation in the single step!!!

"Wild" yet smart idea: Note that the RHS is in the form of "population average". The basic intuitive from statistics is that we can use "sample means" to replace "population average". If you're bold enough -- just randomly pick up ONE single sample and use this value to replace "population average"!

- Heruristic expression of "pure stochastic" SGD:

$$\beta^{k+1} = \beta^k - \eta \nabla_{\beta} l(\beta^k; x^{(r)}, y^{(r)}),$$

where r denotes the index randomly picked during this step.

- (mini-batch SGD, or "standard" SGD):

$$\beta^{k+1} = \beta^k - \eta \frac{1}{n_B} \sum_{k=1}^{n_B} \nabla_{\beta} l(\beta^k; x^{(k)}, y^{(k)}),$$

where n_b denotes the size of mini-batch, and the average is taken over the n_b random samples.

In actual programming, we don't want to generate new random numbers in each step, nor want to "waste" some samples -- we desire all training data can be used during SGD. It is very useful to adopt the "epoch-batch" strategy (or called cyclic rule) through permutation of the data.

Choose initial guess β^0 , step size (learning rate) η ,
batch size n_B , number of inner iterations $M \leq N/n_B$, number of epochs n_E

For epoch $n = 1, 2, \dots, n_E$

β^0 for the current epoch is β^{M+1} for the previous epoch.

Randomly shuffle the training samples.

For $m = 0, 1, 2, \dots, M - 1$

$$\beta^{m+1} = \beta^m - \frac{\eta}{n_B} \sum_{i=1}^{n_B} \nabla_{\beta} l(\beta^m; x^{(m*n_B+i)}, y^{(m*n_B+i)})$$

If the gradient loss of your program is written in a highly vectorized way (support data matrix as input), then you can simply make the data matrix within the mini-batch as the input in each GD update. Below is the example based on our previous binary logistic regression codes.

In practice, you may also find it helpful to adjust the stepsize (learning rate) during the iteration.

```
In [23]: import numpy as np
```

```
class myLogisticRegression_binary():
    """ Logistic Regression classifier -- this only works for the binary case. Here we
    provide the option of SGD in optimization.
    Parameters:
    -----
    learning_rate: float
        The step length that will be taken when following the negative gradient during
        training.
    """
    def __init__(self, learning_rate=.001, opt_method = 'SGD', num_epochs = 50, size_
batch = 20):

        # learning rate can also be in the fit method
        self.learning_rate = learning_rate
        self.opt_method = opt_method
        self.num_epochs = num_epochs
        self.size_batch = size_batch

    def fit(self, data, y, n_iterations = 1000):
        """
        don't forget the document string in methods
        """
        ones = np.ones((data.shape[0],1)) # column of ones
        X = np.concatenate((ones, data), axis = 1) # the augmented matrix, \tilde{X}
        in our lecture
        eta = self.learning_rate

        beta = np.zeros(np.shape(X)[1]) # initialize beta, can be other choices

        if self.opt_method == 'GD':
            for k in range(n_iterations):
                dbeta = self.loss_gradient(beta,X,y) # write another function to compute gradient
                beta = beta - eta * dbeta # the formula of GD
                # this step is optional -- just for inspection purposes
                if k % 500 == 0: # pprint loss every 50 steps
                    print("loss after", k+1, "iterations is: ", self.loss(beta,X,y))

        if self.opt_method == 'SGD':
            N = X.shape[0]
            num_epochs = self.num_epochs
            size_batch = self.size_batch
            num_iter = 0
            for e in range(num_epochs):
                shuffle_index = np.random.permutation(N) # in each epoch, we first re
                shuffle the data to create "randomness"
                for m in range(0,N,size_batch): # m is the starting index of mini-batch
                    i = shuffle_index[m:m+size_batch] # index of samples in the mini-batch
                    dbeta = self.loss_gradient(beta,X[i,:],y[i]) # only use the data
                    in mini-batch to compute gradient. Note the average is taken in the loss_gradient function
                    beta = beta - eta * dbeta # the formula of GD, but this time dbeta is different

                if e % 1 == 0 and num_iter % 50 == 0: # print loss during the training process
                    print("loss after", e+1, "epochs and ", num_iter+1, "iterations is: ", self.loss(beta,X,y))

                num_iter = num_iter + 1 # number of total iterations

            self.coeff = beta

    def predict(self, data):
```

```

ones = np.ones((data.shape[0],1)) # column of ones
X = np.concatenate((ones, data), axis = 1) # the augmented matrix, \tilde{X}
in our lecture
beta = self.coeff # the estimated beta
y_pred = np.round(self.sigmoid(np.dot(X,beta))).astype(int) # >0.5: ->1 else,
->0 -- note that we always use Numpy universal functions when possible
return y_pred

def score(self, data, y_true):
ones = np.ones((data.shape[0],1)) # column of ones
X = np.concatenate((ones, data), axis = 1) # the augmented matrix, \tilde{X}
in our lecture
y_pred = self.predict(data)
acc = np.mean(y_pred == y_true) # number of correct predictions/N
return acc

def sigmoid(self, z):
return 1.0 / (1.0 + np.exp(-z))

def loss(self,beta,X,y):
f_value = self.sigmoid(np.matmul(X,beta))
loss_value = np.log(f_value + 1e-10) * y + (1.0 - y)* np.log(1 - f_value + 1e
-10) # avoid nan issues
return -np.mean(loss_value)

def loss_gradient(self,beta,X,y):
f_value = self.sigmoid(np.matmul(X,beta))
gradient_value = (f_value - y).reshape(-1,1)*X # this is the hardest expressi
on -- check yourself
return np.mean(gradient_value, axis=0)

```

You will find adapting the SGD codes above to multi-class logistic regression is very helpful in doing your final project! (although it's not basic requirement). Here is the very intuitive argument when SGD can boost the algorithms.

Suppose in the training dataset you have $N = 60,000$ samples. With GD, each iteration will cost 60,000 summations. Now consider using SGD. We have the mini-batch size of 30. Then each iteration will cost only 30 sums. For a complete epoch, you have 60,000 sums -- the same with GD, but you have already iterated for 2000 steps!

Of course you may argue that the "quality" of steps in GD is "far better" than SGD. Surely there is the trade-off, but practically [the inferior performance of SGD in convergence does not obscure its super efficiency over GD](https://www.stat.cmu.edu/~ryantibs/convexopt/lectures/stochastic-gd.pdf) (<https://www.stat.cmu.edu/~ryantibs/convexopt/lectures/stochastic-gd.pdf>). In fact, SGD is the de facto optimization method in deep learning. (SGD and BP -- backward propagation to calculate the gradient are the two fundamental cornerstones in deep learning.)

Next, we compare GD and SGD with the UCI ["adult" dataset](https://archive.ics.uci.edu/ml/datasets/adult) (<https://archive.ics.uci.edu/ml/datasets/adult>) to predict income. Note that it is a binary classification problem.

```
In [24]: import pandas as pd
df = pd.read_csv('adult.csv')
df
```

Out[24]:

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship	race	gender	capital-gai
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	Black	Male	
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	Male	
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	Male	
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male	768
4	18	?	103497	Some-college	10	Never-married	?	Own-child	White	Female	
...
48837	27	Private	257302	Assoc-acdm	12	Married-civ-spouse	Tech-support	Wife	White	Female	
48838	40	Private	154374	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband	White	Male	
48839	58	Private	151910	HS-grad	9	Widowed	Adm-clerical	Unmarried	White	Female	
48840	22	Private	201490	HS-grad	9	Never-married	Adm-clerical	Own-child	White	Male	
48841	52	Self-emp-inc	287927	HS-grad	9	Married-civ-spouse	Exec-managerial	Wife	White	Female	1502

48842 rows x 15 columns

```
In [25]: from numpy import nan
df = df.replace('?',nan) #dealing with missing values -- ? in original dataset
df.head()
```

Out[25]:

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship	race	gender	capital-gain	ca
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	Black	Male	0	
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	Male	0	
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	Male	0	
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male	7688	
4	18	NaN	103497	Some-college	10	Never-married	NaN	Own-child	White	Female	0	

```
In [26]: df.dropna(inplace = True) # drop missing values
df
```

Out[26]:

	age	workclass	fnlwgt	education	educational-num	marital-status	occupation	relationship	race	gender	capital gai
0	25	Private	226802	11th	7	Never-married	Machine-op-inspct	Own-child	Black	Male	
1	38	Private	89814	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	Male	
2	28	Local-gov	336951	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	Male	
3	44	Private	160323	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male	768
5	34	Private	198693	10th	6	Never-married	Other-service	Not-in-family	White	Male	
...
48837	27	Private	257302	Assoc-acdm	12	Married-civ-spouse	Tech-support	Wife	White	Female	
48838	40	Private	154374	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband	White	Male	
48839	58	Private	151910	HS-grad	9	Widowed	Adm-clerical	Unmarried	White	Female	
48840	22	Private	201490	HS-grad	9	Never-married	Adm-clerical	Own-child	White	Male	
48841	52	Self-emp-inc	287927	HS-grad	9	Married-civ-spouse	Exec-managerial	Wife	White	Female	1502

45222 rows × 15 columns

```
In [27]: df.drop(columns=['fnlwgt', 'native-country'], inplace=True) # drop some variables we are not interested in
df
```

Out[27]:

	age	workclass	education	educational-num	marital-status	occupation	relationship	race	gender	capital-gain	capital-loss
0	25	Private	11th	7	Never-married	Machine-op-inspct	Own-child	Black	Male	0	
1	38	Private	HS-grad	9	Married-civ-spouse	Farming-fishing	Husband	White	Male	0	
2	28	Local-gov	Assoc-acdm	12	Married-civ-spouse	Protective-serv	Husband	White	Male	0	
3	44	Private	Some-college	10	Married-civ-spouse	Machine-op-inspct	Husband	Black	Male	7688	
5	34	Private	10th	6	Never-married	Other-service	Not-in-family	White	Male	0	
...
48837	27	Private	Assoc-acdm	12	Married-civ-spouse	Tech-support	Wife	White	Female	0	
48838	40	Private	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband	White	Male	0	
48839	58	Private	HS-grad	9	Widowed	Adm-clerical	Unmarried	White	Female	0	
48840	22	Private	HS-grad	9	Never-married	Adm-clerical	Own-child	White	Male	0	
48841	52	Self-emp-inc	HS-grad	9	Married-civ-spouse	Exec-managerial	Wife	White	Female	15024	

45222 rows x 13 columns

```
In [28]: from sklearn.preprocessing import LabelEncoder
df_clean = df.apply(LabelEncoder().fit_transform) # transform the categorical variables into numerical
df_clean
```

Out[28]:

	age	workclass	education	educational-num	marital-status	occupation	relationship	race	gender	capital-gain	capital-loss
0	8	2	1	6	4	6	3	2	1	0	0
1	21	2	11	8	2	4	0	4	1	0	0
2	11	1	7	11	2	10	0	4	1	0	0
3	27	2	15	9	2	6	0	2	1	96	0
5	17	2	0	5	4	7	1	4	1	0	0
...
48837	10	2	7	11	2	12	5	4	0	0	0
48838	23	2	11	8	2	6	0	4	1	0	0
48839	41	2	11	8	6	0	4	4	0	0	0
48840	5	2	11	8	4	0	3	4	1	0	0
48841	35	3	11	8	2	3	5	4	0	110	0

45222 rows × 13 columns

Note that it is not best way to encode the data. Please see other solutions in [kaggle \(https://www.kaggle.com/wenruihu/adult-income-dataset/notebooks\)](https://www.kaggle.com/wenruihu/adult-income-dataset/notebooks).

```
In [29]: y = df_clean['income'].to_numpy()
X = df_clean.drop(columns = 'income').to_numpy()
```

In [30]: X.shape

Out[30]: (45222, 12)

```
In [31]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.1, random_state=42)
```

```
In [32]: from sklearn.linear_model import LogisticRegression
clf = LogisticRegression(random_state=0)
clf.fit(X_train,y_train)
clf.score(X_test,y_test)
```

/Users/cliffzhou/opt/anaconda3/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:940: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG)

Out[32]: 0.8273269953570639

```
In [34]: lg_gd = myLogisticRegression_binary(learning_rate=1e-6, opt_method = 'GD')
lg_sgd = myLogisticRegression_binary(learning_rate=1e-6, opt_method = 'SGD', num_epochs = 15, size_batch = 40)
```



```
In [35]: %%time
lg_gd.fit(X_train,y_train,n_iterations = 15000)

loss after 1 iterations is: 0.6930358550277247
loss after 501 iterations is: 0.6503339171382144
loss after 1001 iterations is: 0.6250322404153786
loss after 1501 iterations is: 0.6091127195017652
loss after 2001 iterations is: 0.5984037857262678
loss after 2501 iterations is: 0.590712724359857
loss after 3001 iterations is: 0.5848586907302407
loss after 3501 iterations is: 0.5801861202580018
loss after 4001 iterations is: 0.5763181444418489
loss after 4501 iterations is: 0.5730292982409765
loss after 5001 iterations is: 0.570178434214311
loss after 5501 iterations is: 0.567672659406349
loss after 6001 iterations is: 0.565447582899603
loss after 6501 iterations is: 0.5634562915787977
loss after 7001 iterations is: 0.5616630677823014
loss after 7501 iterations is: 0.5600397185713462
loss after 8001 iterations is: 0.5585633633136624
loss after 8501 iterations is: 0.5572150485499265
loss after 9001 iterations is: 0.5559788415837271
loss after 9501 iterations is: 0.5548412083490464
loss after 10001 iterations is: 0.5537905657746116
loss after 10501 iterations is: 0.5528169456723574
loss after 11001 iterations is: 0.551911733237817
loss after 11501 iterations is: 0.5510674578925445
loss after 12001 iterations is: 0.5502776225312458
loss after 12501 iterations is: 0.5495365620648746
loss after 13001 iterations is: 0.5488393250200673
loss after 13501 iterations is: 0.548181573717374
loss after 14001 iterations is: 0.5475594996778428
loss after 14501 iterations is: 0.5469697516624197
CPU times: user 2min 14s, sys: 1.88 s, total: 2min 16s
Wall time: 34.5 s
```

```
In [36]: lg_gd.score(X_test,y_test)
```

```
Out[36]: 0.7950475348220207
```

```
In [37]: %%time  
lg_sgd.fit(X_train,y_train)
```

loss after 1 epochs and 1 iterations is: 0.6929943378097103
loss after 1 epochs and 51 iterations is: 0.6879157921439216
loss after 1 epochs and 101 iterations is: 0.6826583387379885
loss after 1 epochs and 151 iterations is: 0.6772727824166075
loss after 1 epochs and 201 iterations is: 0.6726785634817065
loss after 1 epochs and 251 iterations is: 0.6685869015583293
loss after 1 epochs and 301 iterations is: 0.6646473900980252
loss after 1 epochs and 351 iterations is: 0.6611693171938025
loss after 1 epochs and 401 iterations is: 0.657648969777927
loss after 1 epochs and 451 iterations is: 0.6543380353684581
loss after 1 epochs and 501 iterations is: 0.6506396745665634
loss after 1 epochs and 551 iterations is: 0.6475672520727679
loss after 1 epochs and 601 iterations is: 0.6444325842294624
loss after 1 epochs and 651 iterations is: 0.6414499846509609
loss after 1 epochs and 701 iterations is: 0.638868157805757
loss after 1 epochs and 751 iterations is: 0.6359064566095461
loss after 1 epochs and 801 iterations is: 0.6334408621127018
loss after 1 epochs and 851 iterations is: 0.6312278811574829
loss after 1 epochs and 901 iterations is: 0.6290642794693245
loss after 1 epochs and 951 iterations is: 0.6269938964370426
loss after 1 epochs and 1001 iterations is: 0.6250265021072323
loss after 2 epochs and 1051 iterations is: 0.623029627831275
loss after 2 epochs and 1101 iterations is: 0.6211426245735706
loss after 2 epochs and 1151 iterations is: 0.6193467723886272
loss after 2 epochs and 1201 iterations is: 0.6177270519857514
loss after 2 epochs and 1251 iterations is: 0.616107451679061
loss after 2 epochs and 1301 iterations is: 0.6145681630442259
loss after 2 epochs and 1351 iterations is: 0.6130998421319334
loss after 2 epochs and 1401 iterations is: 0.6116220085039598
loss after 2 epochs and 1451 iterations is: 0.6102084881583923
loss after 2 epochs and 1501 iterations is: 0.6090776658809949
loss after 2 epochs and 1551 iterations is: 0.6077928249914799
loss after 2 epochs and 1601 iterations is: 0.6066410887612778
loss after 2 epochs and 1651 iterations is: 0.605489187290032
loss after 2 epochs and 1701 iterations is: 0.6044380017823323
loss after 2 epochs and 1751 iterations is: 0.6033191282399746
loss after 2 epochs and 1801 iterations is: 0.6022723874689804
loss after 2 epochs and 1851 iterations is: 0.6012802739007284
loss after 2 epochs and 1901 iterations is: 0.6002815235883742
loss after 2 epochs and 1951 iterations is: 0.5993572571306436
loss after 2 epochs and 2001 iterations is: 0.5985389048325674
loss after 3 epochs and 2051 iterations is: 0.597531885791725
loss after 3 epochs and 2101 iterations is: 0.5966720291600264
loss after 3 epochs and 2151 iterations is: 0.5957612970107344
loss after 3 epochs and 2201 iterations is: 0.5950766279432791
loss after 3 epochs and 2251 iterations is: 0.5943559676358053
loss after 3 epochs and 2301 iterations is: 0.5935688486792756
loss after 3 epochs and 2351 iterations is: 0.5929201320587859
loss after 3 epochs and 2401 iterations is: 0.5922333698257596
loss after 3 epochs and 2451 iterations is: 0.591532399110286
loss after 3 epochs and 2501 iterations is: 0.5908821436439172
loss after 3 epochs and 2551 iterations is: 0.5902025650949886
loss after 3 epochs and 2601 iterations is: 0.5895305607797755
loss after 3 epochs and 2651 iterations is: 0.5889999803872425
loss after 3 epochs and 2701 iterations is: 0.5883547840173724
loss after 3 epochs and 2751 iterations is: 0.5876780979157352
loss after 3 epochs and 2801 iterations is: 0.5869976541045953
loss after 3 epochs and 2851 iterations is: 0.586354190320411
loss after 3 epochs and 2901 iterations is: 0.5858000785042473
loss after 3 epochs and 2951 iterations is: 0.5853681835590201
loss after 3 epochs and 3001 iterations is: 0.5849082656704451
loss after 3 epochs and 3051 iterations is: 0.5843228468712194
loss after 4 epochs and 3101 iterations is: 0.5838813200076072
loss after 4 epochs and 3151 iterations is: 0.5833637043570097
loss after 4 epochs and 3201 iterations is: 0.5829071213380829
loss after 4 epochs and 3251 iterations is: 0.5824148615973473
loss after 4 epochs and 3301 iterations is: 0.5819581899421266
loss after 4 epochs and 3351 iterations is: 0.5815316681823375
loss after 4 epochs and 3401 iterations is: 0.5811273217174444
loss after 4 epochs and 3451 iterations is: 0.5806500108258908

loss	after	4	epochs	and	3501	iterations	is:	0.5802987640090631
loss	after	4	epochs	and	3551	iterations	is:	0.5799207302452782
loss	after	4	epochs	and	3601	iterations	is:	0.5794817774182126
loss	after	4	epochs	and	3651	iterations	is:	0.5790229197301442
loss	after	4	epochs	and	3701	iterations	is:	0.5786060170770326
loss	after	4	epochs	and	3751	iterations	is:	0.5781610138754161
loss	after	4	epochs	and	3801	iterations	is:	0.5777774569852753
loss	after	4	epochs	and	3851	iterations	is:	0.5773108705786598
loss	after	4	epochs	and	3901	iterations	is:	0.5769441047961352
loss	after	4	epochs	and	3951	iterations	is:	0.5766059315933955
loss	after	4	epochs	and	4001	iterations	is:	0.5763113930219275
loss	after	4	epochs	and	4051	iterations	is:	0.5759388741544418
loss	after	5	epochs	and	4101	iterations	is:	0.5756060993431625
loss	after	5	epochs	and	4151	iterations	is:	0.575246553578359
loss	after	5	epochs	and	4201	iterations	is:	0.5749124798807108
loss	after	5	epochs	and	4251	iterations	is:	0.5745841461780776
loss	after	5	epochs	and	4301	iterations	is:	0.5742173532508857
loss	after	5	epochs	and	4351	iterations	is:	0.5738795327505383
loss	after	5	epochs	and	4401	iterations	is:	0.5735656165110418
loss	after	5	epochs	and	4451	iterations	is:	0.5732648623169405
loss	after	5	epochs	and	4501	iterations	is:	0.5729635942857142
loss	after	5	epochs	and	4551	iterations	is:	0.5726800641595508
loss	after	5	epochs	and	4601	iterations	is:	0.5723953049698364
loss	after	5	epochs	and	4651	iterations	is:	0.5721378941639194
loss	after	5	epochs	and	4701	iterations	is:	0.5718801656792263
loss	after	5	epochs	and	4751	iterations	is:	0.5715722626379177
loss	after	5	epochs	and	4801	iterations	is:	0.5712414660643983
loss	after	5	epochs	and	4851	iterations	is:	0.5709890468205382
loss	after	5	epochs	and	4901	iterations	is:	0.5707011521495051
loss	after	5	epochs	and	4951	iterations	is:	0.5704481533865665
loss	after	5	epochs	and	5001	iterations	is:	0.5701890963151216
loss	after	5	epochs	and	5051	iterations	is:	0.5698944838917015
loss	after	6	epochs	and	5101	iterations	is:	0.5696446485968748
loss	after	6	epochs	and	5151	iterations	is:	0.5694016597723256
loss	after	6	epochs	and	5201	iterations	is:	0.5691625409290855
loss	after	6	epochs	and	5251	iterations	is:	0.5688782405055512
loss	after	6	epochs	and	5301	iterations	is:	0.5685978558901796
loss	after	6	epochs	and	5351	iterations	is:	0.5683716699949697
loss	after	6	epochs	and	5401	iterations	is:	0.5681131998984631
loss	after	6	epochs	and	5451	iterations	is:	0.567877648390682
loss	after	6	epochs	and	5501	iterations	is:	0.5676691209754645
loss	after	6	epochs	and	5551	iterations	is:	0.5674055585744703
loss	after	6	epochs	and	5601	iterations	is:	0.5671662715742651
loss	after	6	epochs	and	5651	iterations	is:	0.5669546220810727
loss	after	6	epochs	and	5701	iterations	is:	0.5667282388169786
loss	after	6	epochs	and	5751	iterations	is:	0.5664845289698955
loss	after	6	epochs	and	5801	iterations	is:	0.5662618263305396
loss	after	6	epochs	and	5851	iterations	is:	0.5660502583056202
loss	after	6	epochs	and	5901	iterations	is:	0.5658594471733267
loss	after	6	epochs	and	5951	iterations	is:	0.5656598505829907
loss	after	6	epochs	and	6001	iterations	is:	0.5654556703079168
loss	after	6	epochs	and	6051	iterations	is:	0.56523953196772
loss	after	6	epochs	and	6101	iterations	is:	0.5650148196411346
loss	after	6	epochs	and	6151	iterations	is:	0.5648148458149036
loss	after	7	epochs	and	6201	iterations	is:	0.5646040718128554
loss	after	7	epochs	and	6251	iterations	is:	0.5643991606267291
loss	after	7	epochs	and	6301	iterations	is:	0.5642087071952897
loss	after	7	epochs	and	6351	iterations	is:	0.5639798886136619
loss	after	7	epochs	and	6401	iterations	is:	0.5637697779922688
loss	after	7	epochs	and	6451	iterations	is:	0.5635938700178585
loss	after	7	epochs	and	6501	iterations	is:	0.5634007419011422
loss	after	7	epochs	and	6551	iterations	is:	0.5632336573333413
loss	after	7	epochs	and	6601	iterations	is:	0.5630560198278177
loss	after	7	epochs	and	6651	iterations	is:	0.5628596155186233
loss	after	7	epochs	and	6701	iterations	is:	0.5626725040341396
loss	after	7	epochs	and	6751	iterations	is:	0.5624882137621866
loss	after	7	epochs	and	6801	iterations	is:	0.5623161020195231
loss	after	7	epochs	and	6851	iterations	is:	0.5621405959168572
loss	after	7	epochs	and	6901	iterations	is:	0.5619819502922073
loss	after	7	epochs	and	6951	iterations	is:	0.561807800793355

loss after 7 epochs and	7001 iterations is:	0.5616302762617381
loss after 7 epochs and	7051 iterations is:	0.5614702327544377
loss after 7 epochs and	7101 iterations is:	0.5613118533738295
loss after 8 epochs and	7151 iterations is:	0.5611911492022321
loss after 8 epochs and	7201 iterations is:	0.5610054312896275
loss after 8 epochs and	7251 iterations is:	0.5608324094471107
loss after 8 epochs and	7301 iterations is:	0.5606755791833046
loss after 8 epochs and	7351 iterations is:	0.5604966305723635
loss after 8 epochs and	7401 iterations is:	0.5603337762624834
loss after 8 epochs and	7451 iterations is:	0.5601687012512644
loss after 8 epochs and	7501 iterations is:	0.56003002480097
loss after 8 epochs and	7551 iterations is:	0.5598876430161809
loss after 8 epochs and	7601 iterations is:	0.5597290059454819
loss after 8 epochs and	7651 iterations is:	0.5595789965911967
loss after 8 epochs and	7701 iterations is:	0.559439370708113
loss after 8 epochs and	7751 iterations is:	0.5592754208065963
loss after 8 epochs and	7801 iterations is:	0.5591224263983647
loss after 8 epochs and	7851 iterations is:	0.5590040310766788
loss after 8 epochs and	7901 iterations is:	0.5588568847916731
loss after 8 epochs and	7951 iterations is:	0.5587250208401735
loss after 8 epochs and	8001 iterations is:	0.5585784612512226
loss after 8 epochs and	8051 iterations is:	0.558417907070479
loss after 8 epochs and	8101 iterations is:	0.5582779557593237
loss after 9 epochs and	8151 iterations is:	0.5581420237672297
loss after 9 epochs and	8201 iterations is:	0.5580445979884451
loss after 9 epochs and	8251 iterations is:	0.5579117251331686
loss after 9 epochs and	8301 iterations is:	0.5577439161296986
loss after 9 epochs and	8351 iterations is:	0.5576019804441052
loss after 9 epochs and	8401 iterations is:	0.5574714227119149
loss after 9 epochs and	8451 iterations is:	0.5573550480444752
loss after 9 epochs and	8501 iterations is:	0.5572138663408357
loss after 9 epochs and	8551 iterations is:	0.5570859635433245
loss after 9 epochs and	8601 iterations is:	0.556956692251072
loss after 9 epochs and	8651 iterations is:	0.5568205125144519
loss after 9 epochs and	8701 iterations is:	0.556692561786039
loss after 9 epochs and	8751 iterations is:	0.5565803103707045
loss after 9 epochs and	8801 iterations is:	0.5564547026608633
loss after 9 epochs and	8851 iterations is:	0.5563081168936544
loss after 9 epochs and	8901 iterations is:	0.5561928351103262
loss after 9 epochs and	8951 iterations is:	0.5560768647185363
loss after 9 epochs and	9001 iterations is:	0.5559662696282236
loss after 9 epochs and	9051 iterations is:	0.5558732332772022
loss after 9 epochs and	9101 iterations is:	0.5557542203984873
loss after 9 epochs and	9151 iterations is:	0.555625593118027
loss after 10 epochs and	9201 iterations is:	0.5554847786984974
loss after 10 epochs and	9251 iterations is:	0.5553888536294705
loss after 10 epochs and	9301 iterations is:	0.5552717784970336
loss after 10 epochs and	9351 iterations is:	0.555128509377872
loss after 10 epochs and	9401 iterations is:	0.5550285606525429
loss after 10 epochs and	9451 iterations is:	0.554923564124486
loss after 10 epochs and	9501 iterations is:	0.55482762399207
loss after 10 epochs and	9551 iterations is:	0.5547358717760673
loss after 10 epochs and	9601 iterations is:	0.5546270101342646
loss after 10 epochs and	9651 iterations is:	0.5545268467536437
loss after 10 epochs and	9701 iterations is:	0.5544299204787582
loss after 10 epochs and	9751 iterations is:	0.5543240574888451
loss after 10 epochs and	9801 iterations is:	0.5542216683885867
loss after 10 epochs and	9851 iterations is:	0.5541239732119122
loss after 10 epochs and	9901 iterations is:	0.5540103214092645
loss after 10 epochs and	9951 iterations is:	0.5539141403632816
loss after 10 epochs and	10001 iterations is:	0.553793756351588
loss after 10 epochs and	10051 iterations is:	0.5536848274180983
loss after 10 epochs and	10101 iterations is:	0.5535975796837489
loss after 10 epochs and	10151 iterations is:	0.5534920651444988
loss after 11 epochs and	10201 iterations is:	0.553381477596243
loss after 11 epochs and	10251 iterations is:	0.5532719602072351
loss after 11 epochs and	10301 iterations is:	0.553174946250866
loss after 11 epochs and	10351 iterations is:	0.5530700187891654
loss after 11 epochs and	10401 iterations is:	0.5529879371619767
loss after 11 epochs and	10451 iterations is:	0.5529028889790654

loss	after	11	epochs	and	10501	iterations	is:	0.5528120075069572
loss	after	11	epochs	and	10551	iterations	is:	0.5527227534626173
loss	after	11	epochs	and	10601	iterations	is:	0.5526103700931397
loss	after	11	epochs	and	10651	iterations	is:	0.5525429379119121
loss	after	11	epochs	and	10701	iterations	is:	0.5524460537628306
loss	after	11	epochs	and	10751	iterations	is:	0.5523581704601671
loss	after	11	epochs	and	10801	iterations	is:	0.5522723133401622
loss	after	11	epochs	and	10851	iterations	is:	0.5521947644454948
loss	after	11	epochs	and	10901	iterations	is:	0.5520825027181196
loss	after	11	epochs	and	10951	iterations	is:	0.5519933166035718
loss	after	11	epochs	and	11001	iterations	is:	0.5519087915779376
loss	after	11	epochs	and	11051	iterations	is:	0.5518136285923861
loss	after	11	epochs	and	11101	iterations	is:	0.5517239891389172
loss	after	11	epochs	and	11151	iterations	is:	0.5516374120142756
loss	after	12	epochs	and	11201	iterations	is:	0.5515658737584687
loss	after	12	epochs	and	11251	iterations	is:	0.5514762012990191
loss	after	12	epochs	and	11301	iterations	is:	0.5513948221996366
loss	after	12	epochs	and	11351	iterations	is:	0.5513189432305279
loss	after	12	epochs	and	11401	iterations	is:	0.551221813621124
loss	after	12	epochs	and	11451	iterations	is:	0.5511252154330428
loss	after	12	epochs	and	11501	iterations	is:	0.5510311755495879
loss	after	12	epochs	and	11551	iterations	is:	0.5509381113325423
loss	after	12	epochs	and	11601	iterations	is:	0.5508610809845267
loss	after	12	epochs	and	11651	iterations	is:	0.5507766641400444
loss	after	12	epochs	and	11701	iterations	is:	0.5506998417700848
loss	after	12	epochs	and	11751	iterations	is:	0.550638054602862
loss	after	12	epochs	and	11801	iterations	is:	0.5505767388998226
loss	after	12	epochs	and	11851	iterations	is:	0.550497291395151
loss	after	12	epochs	and	11901	iterations	is:	0.5504046964469879
loss	after	12	epochs	and	11951	iterations	is:	0.550331781834684
loss	after	12	epochs	and	12001	iterations	is:	0.550260432566219
loss	after	12	epochs	and	12051	iterations	is:	0.5501907117369469
loss	after	12	epochs	and	12101	iterations	is:	0.5501195584752477
loss	after	12	epochs	and	12151	iterations	is:	0.5500477900330109
loss	after	12	epochs	and	12201	iterations	is:	0.5499768139505086
loss	after	13	epochs	and	12251	iterations	is:	0.5499109132983397
loss	after	13	epochs	and	12301	iterations	is:	0.5498309833565308
loss	after	13	epochs	and	12351	iterations	is:	0.549771388818261
loss	after	13	epochs	and	12401	iterations	is:	0.5496965511927396
loss	after	13	epochs	and	12451	iterations	is:	0.5496140938400194
loss	after	13	epochs	and	12501	iterations	is:	0.549539115576209
loss	after	13	epochs	and	12551	iterations	is:	0.5494855492956467
loss	after	13	epochs	and	12601	iterations	is:	0.5494094085657938
loss	after	13	epochs	and	12651	iterations	is:	0.5493316028948573
loss	after	13	epochs	and	12701	iterations	is:	0.5492563140350402
loss	after	13	epochs	and	12751	iterations	is:	0.5491979481772212
loss	after	13	epochs	and	12801	iterations	is:	0.5491447078250461
loss	after	13	epochs	and	12851	iterations	is:	0.5490756472727727
loss	after	13	epochs	and	12901	iterations	is:	0.5489993369554883
loss	after	13	epochs	and	12951	iterations	is:	0.5489259736410982
loss	after	13	epochs	and	13001	iterations	is:	0.5488520237105091
loss	after	13	epochs	and	13051	iterations	is:	0.5487764619075166
loss	after	13	epochs	and	13101	iterations	is:	0.5487068900319433
loss	after	13	epochs	and	13151	iterations	is:	0.5486429723126187
loss	after	13	epochs	and	13201	iterations	is:	0.548576954914066
loss	after	14	epochs	and	13251	iterations	is:	0.5485066144257112
loss	after	14	epochs	and	13301	iterations	is:	0.5484402220055412
loss	after	14	epochs	and	13351	iterations	is:	0.5483861516730834
loss	after	14	epochs	and	13401	iterations	is:	0.5483093245103873
loss	after	14	epochs	and	13451	iterations	is:	0.5482484298640372
loss	after	14	epochs	and	13501	iterations	is:	0.5481943437165917
loss	after	14	epochs	and	13551	iterations	is:	0.5481246750116509
loss	after	14	epochs	and	13601	iterations	is:	0.5480607208175462
loss	after	14	epochs	and	13651	iterations	is:	0.5479883239196417
loss	after	14	epochs	and	13701	iterations	is:	0.547932114906443
loss	after	14	epochs	and	13751	iterations	is:	0.5478832774457331
loss	after	14	epochs	and	13801	iterations	is:	0.5478191196664128
loss	after	14	epochs	and	13851	iterations	is:	0.5477627692925748
loss	after	14	epochs	and	13901	iterations	is:	0.5476951793177156
loss	after	14	epochs	and	13951	iterations	is:	0.547636942391832

```
loss after 14 epochs and 14001 iterations is: 0.5475727038831864
loss after 14 epochs and 14051 iterations is: 0.5475029962872622
loss after 14 epochs and 14101 iterations is: 0.5474522011473978
loss after 14 epochs and 14151 iterations is: 0.5473888639126175
loss after 14 epochs and 14201 iterations is: 0.5473191568466856
loss after 14 epochs and 14251 iterations is: 0.5472595767863389
loss after 15 epochs and 14301 iterations is: 0.5471959816915696
loss after 15 epochs and 14351 iterations is: 0.547139600171737
loss after 15 epochs and 14401 iterations is: 0.547070942183188
loss after 15 epochs and 14451 iterations is: 0.5470157722029524
loss after 15 epochs and 14501 iterations is: 0.5469641316728932
loss after 15 epochs and 14551 iterations is: 0.5469066325243004
loss after 15 epochs and 14601 iterations is: 0.5468556034655901
loss after 15 epochs and 14651 iterations is: 0.5467853514796428
loss after 15 epochs and 14701 iterations is: 0.5467286750489505
loss after 15 epochs and 14751 iterations is: 0.5466789227751985
loss after 15 epochs and 14801 iterations is: 0.5466193858166553
loss after 15 epochs and 14851 iterations is: 0.5465650457604271
loss after 15 epochs and 14901 iterations is: 0.5465101745326028
loss after 15 epochs and 14951 iterations is: 0.5464607661008346
loss after 15 epochs and 15001 iterations is: 0.5464136330326027
loss after 15 epochs and 15051 iterations is: 0.5463591185069542
loss after 15 epochs and 15101 iterations is: 0.5463038550222042
loss after 15 epochs and 15151 iterations is: 0.5462387689962773
loss after 15 epochs and 15201 iterations is: 0.5461904968259284
loss after 15 epochs and 15251 iterations is: 0.5461404214621863
CPU times: user 6.8 s, sys: 631 ms, total: 7.43 s
Wall time: 1.86 s
```

```
In [38]: lg_sgd.score(X_test,y_test)
```

```
Out[38]: 0.7952686270174663
```